

Reducing False Alarms in Edge Driver Monitoring Systems through Negative Frame Curation and Hard Negative Mining

Author names are omitted for double-blind review

Abstract—Frequent false alarms in edge-deployed driver monitoring systems erode driver trust, waste on-device compute, and raise thermal load. This paper evaluates whether curating negative training frames reduces false alarms, and whether replacing randomly sampled negatives with hard-mined false positives suppresses them further. Four lightweight You Only Look Once (YOLO) detectors were trained and benchmarked on a 15,723-frame dataset extracted from the Driver Monitoring Dataset across four localized visual cues, using a dedicated graphics processing unit. Models were evaluated across negative-frame ratios from 0% to 80%, followed by hard-negative mining trials and a latency benchmark measuring per-frame runtime against the standard 30 ms real-time in-cabin deadline, giving an edge-proxy compute baseline ahead of low-power embedded deployment. Increasing negative exposure consistently reduced false positives across all detectors, achieving a 6.4× to 25.7× reduction at 80% while detection performance stayed stable. At the 40% reference ratio, hard negative mining cut false positives a further 33.3% to 67.8% versus random sampling, with the gain varying by model. Runtime profiling showed sub-30 ms inference latency at batch size one without altering model architecture or complexity. Negative frame curation and hard negative mining together offer a computationally inexpensive way to improve the operational reliability of lightweight edge Driver Monitoring Systems.

Index Terms—Driver monitoring systems, object detection, false alarm suppression, hard-negative mining, edge computing.

I. INTRODUCTION

Beginning 7 July 2024, European Union (EU) regulations required Advanced Driver Distraction Warning (ADDW) systems in newly type-approved vehicles, extending to all newly registered vehicles from 7 July 2026 [1]. Early vision-based Driver Monitoring Systems (DMS) estimated driver states via spatio-temporal action recognition [2], [8], but these three-dimensional convolutional architectures demand excessive memory and compute time, making them unsuitable for low-power automotive microcontrollers [3], [8]. Modern edge DMS pipelines instead formulate monitoring as frame-by-frame object detection [9], [11], deploying lightweight models such as the You Only Look Once (YOLO) detector family to identify localized behavioral cues including phone use, drinking, and eye closure [2], [10]. In-cabin edge hardware must run these models under strict limits: inference latency under 30 ms [4], and power draw below 10–15 W to avoid overheating [3], [6].

Lightweight detectors meet these power and latency targets, but existing DMS datasets consist almost entirely of positive distraction events [9], [11]. Lacking temporal context, unreg-

ularized detectors trigger frequent false alarms on steering wheels, seat belts, and clothing folds [8]. Naturalistic on-road studies show that distracting secondary-task engagement occupies only a small share of driving time—one widely cited video-based study measured such engagement, excluding passenger conversation, at just 14.5% [18]—so cue-free driving dominates on-road operation. Our evaluation stream standardizes background prevalence at $p_{\text{neg}} = 0.809$ (matching our held-out test split) to benchmark false-alarm suppression under controlled in-cabin conditions. Frequent false alarms cause alert fatigue, leading drivers to ignore, mute, or bypass safety systems entirely [5], [13], and also trigger downstream tasks such as facial landmark tracking, Controller Area Network (CAN) bus logging, and video recording [3], [4], creating processing spikes that cause thermal throttling on embedded chips [6]. Adding negative (background-only) frames during training reduces false alarms, but the impact of negative-sampling ratios and hard-negative mining remains under-evaluated in edge DMS applications [16], [17].

To address these challenges, this paper provides three main contributions:

- **Negative Sampling Benchmarks:** A systematic evaluation of negative-to-positive frame ratios (0%–80%) across four lightweight edge-oriented detectors (YOLO11n, YOLO26n, YOLO12n, YOLOv10n) [10], [12].
- **Hard-Negative Mining Analysis:** An experimental evaluation showing targeted hard-negative mining lowers the false alarm rate (FAR) by up to 67.8% without reducing cue detection accuracy [14].
- **Operational Nuisance Alert Modeling:** A formulation bridging frame-level false alarm rates to hourly nuisance alerts ($\mathcal{A}_h$) across streaming regimes (5, 15, 30 FPS), paired with latency profiling against the 30 ms deadline to assess operational reliability and cascading thermal overhead.

Specifically, this study answers two central research questions (RQs):

- **RQ1:** What ratio of negative-to-positive training frames minimizes false alarms without reducing cue detection sensitivity, and is this ratio consistent across different lightweight architectures?
- **RQ2:** Does replacing random negative frames with hard negatives (prior false positives) provide better false-alarm suppression in resource-constrained models?

II. RELATED WORK

A. Vision-Based Driver Monitoring and Cue Detection

Vision-based DMS historically framed driver state estimation as action recognition [2], [8], using convolutional neural networks (CNNs) to classify video clips into coarse categories such as safe driving, talking, or drinking [2], [8]. Multimodal benchmarks, including the Driver Monitoring Dataset (DMD) [2] and Drive&Act [7], established standardized evaluation protocols across red-green-blue (RGB), depth, and infrared streams. However, 3D CNNs and long short-term memory (LSTM) networks impose latency and memory demands exceeding automotive microcontroller limits [3], [8].

To satisfy real-time constraints, modern edge DMS pipelines instead reframe monitoring as frame-by-frame object detection using lightweight YOLO models [9], [11]. Recent works propose customized architectures: YOLO-Drive adds state-space blocks and attention mechanisms [9], P-YOLOv8 pairs a pretrained YOLOv8 backbone with a lightweight classification head for TinyML deployment [11], and Chaman *et al.* benchmark compact YOLOv12 models [10]. Yet these detectors train almost entirely on positive distraction intervals [9], [11]. Without background regularization, single-frame models lack temporal context and frequently misclassify steering wheels, seat belts, and clothing folds as distraction cues [8]. Rather than designing bespoke architectures, this work offers a data-centric evaluation of negative-frame curation and hard-negative mining to suppress false alarms across standard YOLO detectors.

B. Automotive Edge Computing and Cascading Thermal Overhead

Automotive edge platforms must sustain sub-30 ms inference within a 10–15 W power budget under passive cooling [3], [4]. Embedded DMS studies typically evaluate frame rates and nominal power under static workloads [3], yet deployed pipelines trigger secondary tasks on each cue detection, including facial landmark alignment, gaze tracking, CAN bus logging, and video ring-buffer encoding [3], [4].

When an unregularized detector produces frequent false alarms, these secondary tasks execute repeatedly. Benoit-Cattin *et al.* showed that sustained processing bursts cause thermal throttling in embedded vision processors, triggering dynamic frequency scaling that cuts inference throughput by over 40% [6]. Existing DMS benchmarks evaluate detectors in isolation, overlooking how frame-level false positives compound into hourly nuisance alerts ($\mathcal{A}_h$) that induce thermal stress.

C. Alert Fatigue, Automation Disuse, and Regulatory Mandates

European Union (EU) vehicle-safety regulations require reliable distraction detection and prompt warning escalation [1], yet persistent false alarms compromise safety-system efficacy. In human-automation interaction theory, Parasuraman and Riley established that false alarms are the primary cause of automation *disuse*, the intentional muting, ignoring, or

disabling of safety systems by operators [13]. Repeated false alerts during routine driving induce alert fatigue and erode driver trust, motivating motorists to circumvent monitoring hardware with aftermarket defeat devices [5], [13]. Suppressing false alarms is therefore essential for driver acceptance and regulatory compliance.

D. Negative-Frame Curation and Hard-Negative Mining

In object detection, false-positive suppression typically relies on loss reweighting such as Focal Loss [15], or candidate resampling such as Online Hard Example Mining (OHEM) [14]. However, online anchor-level mining like OHEM is computationally impractical for lightweight single-stage architectures, and loss reweighting cannot compensate when training sets lack background negatives entirely [15]. Offline image-level bootstrapping, harvesting false-positive frames from unannotated streams to augment later training passes, offers a computationally viable alternative for compact detectors.

Practitioner heuristics often recommend adding 0% to 10% background frames during training [16], but this guideline targets general object detection benchmarks and remains unvalidated in automotive environments, where naturalistic on-road evidence indicates that distracting secondary-task engagement occupies only a small share of driving time [18], and standard benchmarks under-sample these baseline, cue-free periods. Uniform random negative sampling also yields diminishing returns, since simple background scenes generate minimal gradient feedback [14], [17]. While video-based mining techniques harvest temporal variations as hard negatives in generic domains [17], prior work has not systematically evaluated negative ratios from 0% to 80% across lightweight YOLO detectors, nor measured hard-negative mining’s impact on operational nuisance alert rates ($\mathcal{A}_h$).

III. METHODOLOGY

This section details the models, dataset, split protocols, metrics, and experimental procedures addressing two research questions:

- **RQ1 (Negative Sampling Ratio):** What ratio of negative-to-positive training frames minimizes background false alarms without degrading cue sensitivity, and is this threshold consistent across lightweight architectures?
- **RQ2 (Hard-Negative Mining):** Does replacing random background frames with hard negatives (prior false positives) give superior suppression at matched training set size?

Sections III-A–III-D define architectures, dataset curation, splits, and metrics for RQ1. Sections III-E–III-F detail the mining protocol and deployment benchmarks for RQ2.

A. Model Architectures and Training Configuration

We benchmark four lightweight You Only Look Once (YOLO) detectors under 3 million parameters: YOLO11n [20], YOLO26n [22], YOLO12n [21], and YOLOv10n [19]. Table I

summarizes each model’s parameter count, giga floating-point operations (GFLOPs), and input resolution.

TABLE I
EVALUATED YOLO ARCHITECTURES

Model	Params (M)	GFLOPs	Input Size
YOLO11n	2.6	6.5	640×640
YOLO26n	2.6	6.2	640×640
YOLO12n	2.6	7.5 [†]	640×640
YOLOv10n	2.3	6.6	640×640

[†] Profiled via `thop` at FP32; vendor specification reports 6.5 GFLOPs.

TABLE II
TRAINING HARDWARE AND HYPERPARAMETERS

Parameter	Value
GPU	NVIDIA RTX 4060 (8GB VRAM)
Pretrained weights	COCO
Epochs	100
Batch size	16
Optimizer	AdamW
Initial LR	0.00125
Weight decay	0.0005
Warmup epochs	3.0
LR schedule	Linear decay
Input size	640×640

All models are initialized from Common Objects in Context (COCO) pretrained checkpoints and trained via Ultralytics using the configuration in Table II on an NVIDIA RTX 4060 graphics processing unit (GPU) with 8GB VRAM. Dataloader workers and PyTorch random generators were seeded ($\text{seed} \in \{42, 43\}$) under `deterministic=True` for reproducible batch sequencing. Augmentation includes horizontal flipping ($p = 0.5$), hue-saturation-value (HSV) shifts, translation ($\pm 10\%$), random erasing ($p = 0.4$), and mosaic augmentation (disabled for the final 10 epochs via `close_mosaic=10`). Geometric distortions (rotation, shear, perspective warp) are omitted to preserve in-cabin camera geometry.

B. Dataset Curation and Annotation Protocol

We extract 15,723 face-crop frames at 1 frame per second (FPS) from 81 red-green-blue (RGB) recordings across 14 drivers in the Driver Monitoring Dataset (DMD) [2]. The curated set contains 3,001 positive cue frames (19.1%) and 12,722 negative frames (80.9%), giving a background-heavy stream that approximates the operational rarity of distraction events. Negative frames contain no bounding boxes; positive frames contain exactly one localized bounding box identifying one of four driver cues: yawning, hand-over-mouth, drinking, or phone use (Fig. 1(a)).

Annotations were generated by a single annotator under three rules:

- **Mutual Exclusion:** If a driver covers a yawn with a hand, the behavior is labeled exclusively as hand-over-mouth.

- **Visible Anatomy Only:** Bounding boxes enclose only directly visible cue interactions without extrapolating occluded joints.
- **Scale Invariance:** No minimum bounding-box pixel dimension is imposed at the 640×640 resolution.

C. Training Splits, Negative Ratios, and Seed Protocol

We apply an 80/10/10 stratified split across the 3,001 positive frames, allocating 2,401 positive frames to training and 300 each to validation and testing.

TABLE III
DATASET SPLITS AND FRAME DISTRIBUTIONS

Split	Frame Counts		
	Positive	Negative	Total
0%	2,401	0	2,401
20%	2,401	600	3,001
40%	2,401	1,600	4,001
60%	2,401	3,602	6,003
80%	2,401	9,604	12,005
Val [†]	300	1,272	1,572
Test [†]	300	1,272	1,572

[†] Held-out evaluation splits feature an 80.9% negative ratio.

TABLE IV
RANDOM SEED PROTOCOL BY NEGATIVE RATIO

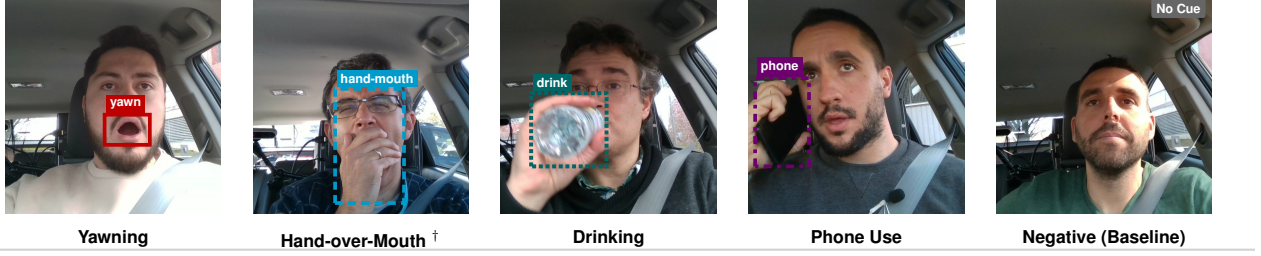
Negative Ratio	Seeds Used
0%, 20%, 40%	42, 43
60%, 80%	42

For RQ1, we build five training configurations by augmenting the 2,401 positive frames with negative-frame proportions from 0% to 80% (Table III). Both validation and test sets keep a fixed 80.9% negative ratio (1,272 negative and 300 positive frames, 1,572 total each) to benchmark false-positive resistance under a standardized background-heavy distribution. To account for initialization and batch-shuffling variance, configurations under sparse negative supervision (0%, 20%, 40%) were trained across two random seeds ($\text{seed} \in \{42, 43\}$) and reported as dual-seed means. Given the compute budget for the 20-configuration matrix, the 60% and 80% sweeps used a single baseline seed ($\text{seed} = 42$), per Table IV. The monotonic reduction across all four architectures corroborates the stability of the 80% ratio, though future work should extend multi-seed coverage uniformly across all ratios.

D. Evaluation Metrics and False-Alarm Quantification

Models are evaluated on detection precision, recall, F1 score, mean Average Precision (mAP), false alarm rate (FAR), and false positives per 1,000 frames (FP/1k). Detections undergo non-maximum suppression (NMS) at an IoU threshold of 0.70 and an operational confidence threshold of $\tau = 0.25$. While raising τ post-hoc lowers false alarms, it degrades recall on subtle cues (e.g., partial eye closure), necessitating background negative regularization at fixed sensitivity.

(a) Localized Behavioral Cue Annotations vs. Baseline Negative



(b) Training Pipeline and Closed-Loop Hard-Negative Mining



Fig. 1. Overview of the training and deployment pipeline: (a) Target behavioral cues (640×640 pixels) with an unannotated negative frame (baseline driving). A covered yawn is labeled hand-over-mouth ($\dagger$); line styles and colors provide dual encoding for greyscale clarity. (b) The pipeline sweeps negative-frame ratios and harvests false positives (FPs) to retrain the model, reducing projected nuisance alert rates ($\mathcal{A}_h$) for in-cabin driver monitoring systems (DMS); the final stage reflects desktop-GPU latency profiling used as a compute proxy, not validated automotive SoC hardware.

A prediction is scored as a True Positive (TP) if it predicts the correct cue class with Intersection over Union (IoU) ≥ 0.50 against a ground-truth box; otherwise it is a False Positive (FP). Ground-truth boxes with no matching detection are False Negatives (FN). Precision, recall, and optimal F1 score are defined as:

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}, \quad \text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}, \quad (1)$$

$$\text{F1} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}. \quad (2)$$

Localization quality is measured using mAP_{50} across the $C = 4$ driver-cue classes:

$$\text{mAP}_{50} = \frac{1}{C} \sum_{c=1}^C \text{AP}_c^{50}, \quad (3)$$

where AP_c^{50} is the Average Precision for class c at IoU = 0.50. We also report $\text{mAP}_{50:95}$, averaged across ten IoU thresholds from 0.50 to 0.95 in 0.05 increments.

To quantify false alerts on background scenes, let FP_{neg} denote detections with confidence $\geq \tau$ ($\tau = 0.25$) produced on the $N_{\text{neg}} = 1,272$ unannotated test frames ($\mathcal{D}_{\text{eval}}^{\text{neg}}$). Per-frame FAR and scaled operational FP/1k are defined as:

$$\text{FAR} = \frac{FP_{\text{neg}}}{N_{\text{neg}}}, \quad \text{FP/1k} = 1,000 \times \text{FAR}. \quad (4)$$

E. Hard-Negative Mining Protocol

To address RQ2, we implement the closed-loop hard-negative mining protocol shown in Fig. 1(b). We first train

a baseline model f_{θ_0} on the 0% negative baseline (2,401 positive frames). Lacking background regularization, f_{θ_0} readily triggers on benign cabin features, making it an effective filter for discovering informative non-target patterns.

We run f_{θ_0} in batched inference over unselected training background frames ($\mathcal{U}_{\text{neg}}$, strictly isolated from validation and test sets). For frame $X \in \mathcal{U}_{\text{neg}}$, let $\text{conf}(X)$ denote the maximum detection confidence. Frames triggering false detections with $\text{conf}(X) \geq \tau$ ($\tau = 0.25$) are harvested as hard negatives:

$$\mathcal{H}_{\text{mined}} = \{X \in \mathcal{U}_{\text{neg}} \mid \text{conf}(X) \geq \tau\}. \quad (5)$$

Because the harvested hard-negative count varies across architectures (Fig. 2), we fix the total negative quota to $N_{\text{target}} = 1,600$, matching the 40% reference curation ratio (4,001 total frames, Table III). Frames in $\mathcal{H}_{\text{mined}}$ are ranked by descending confidence, and remaining slots are backfilled via uniform random sampling (seed = 42) from $\mathcal{U}_{\text{neg}} \setminus \mathcal{H}_{\text{mined}}$. Drawing both 40% variants from the same 10,178-frame pool at fixed cardinality ($N = 1,600$) isolates sample informativeness from training set volume. Detectors are retrained from COCO checkpoints across two seeds (seed $\in \{42, 43\}$) using the baseline training protocol.

F. Real-Time Deployment and Runtime Overhead Evaluation

To quantify operational impact, we project empirical FAR into expected hourly nuisance alerts, $\mathcal{A}_h$. For camera streaming rate r FPS, one hour of driving (3,600 s), and background

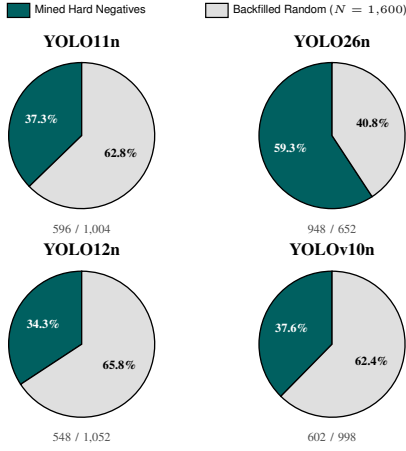


Fig. 2. Proportion of mined false positives versus backfilled random negatives per model (40% curation ratio, total $N = 1,600$). Subtitles report exact mined / backfilled frame counts.

driving prevalence $p_{\text{neg}} \approx 0.809$, assuming independent per-frame inference:

$$\mathcal{A}_h = 3600 \times r \times p_{\text{neg}} \times \text{FAR}. \quad (6)$$

Real-world false detections from static cabin features (such as seatbelt brackets or steering wheel seams) often occur in temporally correlated bursts rather than independent Bernoulli trials, so (6) gives an unmitigated upper-bound rate for comparing raw detector false-alarm resistance. We evaluate $\mathcal{A}_h$ across three streaming rates, $r \in \{5, 15, 30\}$ FPS, corresponding to low-power sampling, typical embedded vision capture, and full real-time streaming.

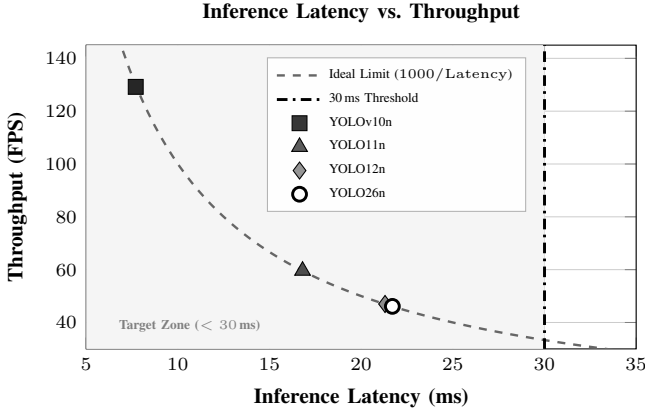


Fig. 3. Inference latency versus throughput across the evaluated YOLO models. Each model uses a distinct marker shape for greyscale clarity. The dashed vertical line marks the 30 ms real-time deadline.

Runtime efficiency is profiled on an NVIDIA RTX 4060 GPU using single-image batches (640×640 , FP32) averaged over 200 iterations after 50 warm-up runs. Dedicated automotive system-on-chip (SoC) hardware with passive thermal limits (10–15 W) is the ultimate target deployment, but this GPU benchmark establishes an execution latency baseline. As

shown in Fig. 3, all four models run well within the 30 ms real-time deadline (7.7–21.7 ms per frame, 46.1–129.2 FPS) with peak memory under 1.5 GB VRAM, confirming that architectural complexity does not violate frame-rate requirements.

IV. EXPERIMENTAL RESULTS

A. Negative-to-Positive Ratio Sweep

To address RQ1, we evaluate each detector across five negative-to-positive training ratios, 0%, 20%, 40%, 60%, and 80% (0:1 to 4.00:1; Table III). Fig. 4 presents detection accuracy and false-positive rates on the held-out test set (1,272 negative frames); full per-configuration metrics are reported in Table V.

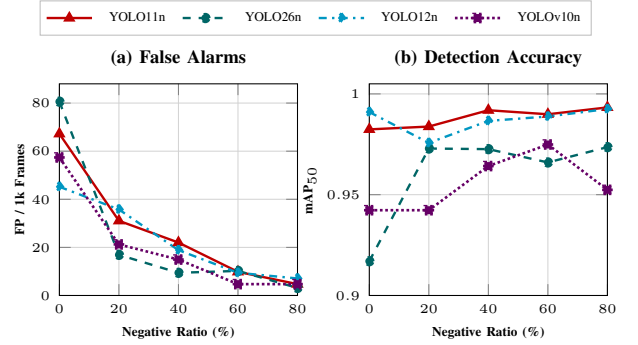


Fig. 4. Impact of negative-to-positive curation ratios (RQ1): (a) false alarms per 1,000 frames (FP/1k) across curation ratios; (b) core detection accuracy (mAP₅₀). Distinct marker shapes and line styles provide visual encoding for greyscale readability.

TABLE V
NEGATIVE-FRAME RATIO SWEEP ACROSS ARCHITECTURES (RQ1)

Det.	r	mAP ₅₀	mAP _{50:95}	P	R	FP/1k
YOLO11n	0%	0.9824	0.7570	0.8576	0.9787	67.22
	20%	0.9838	0.7524	0.9283	0.9814	31.05
	40%	0.9919	0.7711	0.9308	0.9831	22.01
	60%	0.9899	0.7723	0.9468	0.9806	9.83
	80%	0.9933	0.7606	0.9644	0.9766	4.72
YOLO26n	0%	0.9170	0.7086	0.8557	0.9105	80.59
	20%	0.9729	0.7546	0.9306	0.9763	16.90
	40%	0.9726	0.7388	0.9224	0.9519	9.43
	60%	0.9660	0.7594	0.9587	0.9086	10.22
	80%	0.9736	0.7671	0.9217	0.9794	3.14
YOLO12n	0%	0.9910	0.7496	0.9228	0.9823	45.20
	20%	0.9757	0.7449	0.8954	0.9863	35.77
	40%	0.9866	0.7659	0.9199	0.9949	18.86
	60%	0.9888	0.7766	0.9348	0.9960	9.43
	80%	0.9925	0.7529	0.9621	0.9767	7.08
YOLOv10n	0%	0.9423	0.7287	0.9065	0.8831	57.39
	20%	0.9423	0.7270	0.8833	0.9400	21.23
	40%	0.9642	0.7568	0.8813	0.9788	14.94
	60%	0.9749	0.7525	0.9221	0.9758	4.72
	80%	0.9524	0.7384	0.9010	0.9024	4.72

Dual-seed means (seed $\in \{42, 43\}$) for $r \leq 40\%$; single seed (seed = 42) otherwise. FP/1k is computed at the fixed operational threshold $\tau = 0.25$ (NMS IoU 0.70) over 1,272 negative test frames. Precision (P) and recall (R) are reported at the confidence threshold maximizing F_1 on the validation split for each configuration, which may differ from $\tau = 0.25$; the two are distinct operating points and should not be read as arising from the same confusion matrix. Bold indicates best performance per architecture.

Convergence at High Ratios: Detectors trained without negative frames (0%) show high false-alarm rates of 45.20 to 80.59 FP/1k (Fig. 4). Raising negative frames to 80% reduces false alarms across all models to a narrow 3.14–7.08 FP/1k range, a $6.4\times$ to $25.7\times$ reduction. Negative ratio strongly inversely correlates with FP/1k (Spearman $\rho = -0.937, p < 10^{-4}$ across all 20 configurations, reported descriptively without repeated-measures cluster corrections).

Cue Sensitivity and Localization Stability: Precision improves significantly with negative ratio (Spearman $\rho = +0.546, p < 0.05$), directly reflecting false-alarm suppression. Recall stays uncorrelated with negative ratio (Spearman $\rho = +0.012, p > 0.50$), confirming negative frames do not cause under-detection. Localization accuracy ($\text{mAP}_{50:95}$) stays within a narrow band across all configurations (0.7086–0.7766), showing no systematic degradation with increasing negative exposure.

Intermediate Trajectory Differences: Although all models converge at 80% negatives, trajectories differ between 20% and 60%. YOLO26n drops sharply to 16.90 FP/1k at 20%, plateauing early. YOLO11n and YOLO12n decline more gradually, requiring ratios between 60% and 80%, whereas YOLOv10n suppresses early by 40% (14.94 FP/1k) before reaching its 4.72 FP/1k floor at 60%. Fig. 5 consolidates these trends on a shared logarithmic scale alongside the projected hourly nuisance-alert axis, overlaying the 40% hard-negative mining results (Section IV-B) against random sampling at the same ratio.

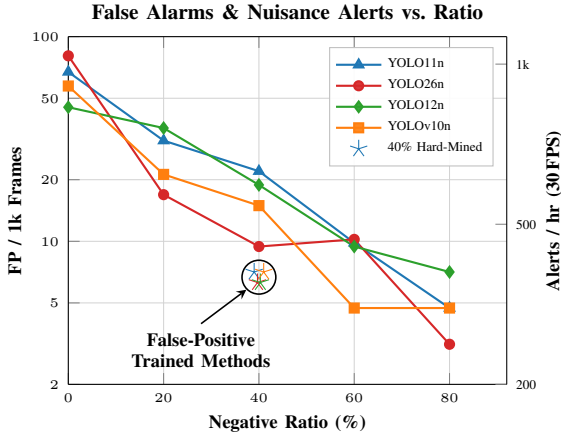


Fig. 5. Measured false-positive rates (left axis) and projected hourly alerts at 30 FPS (right axis; via (6)) versus negative ratio. Both axes share synchronized log scales due to the linear relation ($\mathcal{A}_h = 87.4 \times \text{FP/1k}$). The circled cluster highlights 40% hard-negative mining results; star markers are jittered (39.0–41.0) for visual separation only, all four represent the identical 40% ratio.

B. Hard-Negative Mining vs. Random Negatives

To address RQ2, we evaluate whether hard negatives improve suppression over random background frames at matched dataset volume. Following Section III-E, we harvest false positives from the 0% baseline detector ($\tau = 0.25$) and backfill with random negatives to reach 1,600 negative frames (4,001

total, matching the 40% ratio). Fig. 6 presents two-seed sample means (seed $\in \{42, 43\}$) comparing random sampling with hard-negative mining.

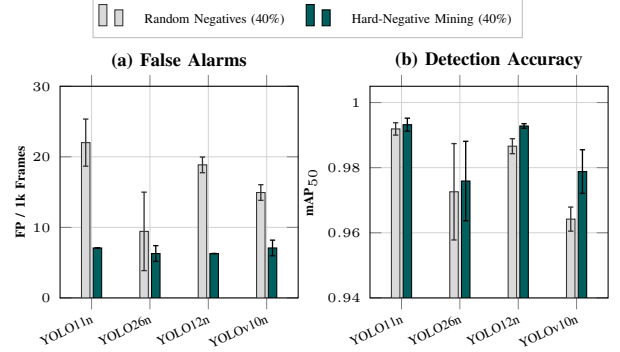


Fig. 6. Hard-negative mining versus random negative sampling ($r_{\text{ref}} = 40\%$, RQ2): (a) false alarms per 1,000 frames (FP/1k); (b) core detection accuracy (mAP_{50}). Bars report dual-seed sample means with ± 1 standard deviation error bars.

Architectural Variation in Suppression: Hard mining reduces false alarms across all four architectures, but the magnitude depends on model design (Fig. 6):

- *Attention and Dual-Assignment Models:* YOLO11n, YOLO12n, and YOLOv10n achieve substantial reductions: -67.8% for YOLO11n ($22.01 \rightarrow 7.08$ FP/1k), -66.7% for YOLO12n ($18.86 \rightarrow 6.29$ FP/1k), and -52.6% for YOLOv10n ($14.94 \rightarrow 7.08$ FP/1k). Localization quality also improves, with $\text{mAP}_{50:95}$ up to $+0.018$ higher for YOLO12n.
- *Purely Convolutional Models:* YOLO26n shows a smaller -33.3% reduction ($9.43 \rightarrow 6.29$ FP/1k), reflecting diminishing returns from targeted mining over random background frames.

With dual-seed training ($n = 2$), these comparisons reflect empirical sample mean differences rather than formal hypothesis tests; given the variance in some configurations (e.g., YOLO26n random-40% standard deviation of ± 5.56 FP/1k), larger multi-seed batteries are needed to establish statistical significance.

C. Nuisance Alert Projection

To evaluate operational reliability, we project measured false-alarm rates into expected hourly nuisance alerts ($\mathcal{A}_h$) using (6). On-road naturalistic driving research indicates distracting secondary-task engagement occupies only a minority of driving time [18]; we nonetheless use our held-out test prevalence of $p_{\text{neg}} = 0.809$ for consistency with our test partition. Alert rates scale directly across operational streaming rates: 5 FPS (low-power sampling), 15 FPS (standard embedded vision), and 30 FPS (full real-time streaming, mapped on the right axis of Fig. 5).

Alert Fatigue and Automation Disuse: Models trained without negative frames (0%) trigger impractically high alert frequencies (Fig. 5). At 30 FPS, unregularized detectors produce 3,949 to 7,041 alerts/hr (1–2 false alerts per second); even

at 5 FPS they generate 658 to 1,174 alerts/hr. This frequency exceeds driver tolerance in human-automation interaction theory, directly inducing alert fatigue and automation disuse [5], [13].

Cascading Edge Compute and Thermal Load: Continuous false alarms repeatedly trigger downstream facial tracking, logging, and video encoding routines. The resulting processing spikes can induce thermal throttling and cut edge throughput by over 40% [3], [6].

Mitigation via Negative Curation: Curating negative training frames lowers projected alert rates by up to 96.1%. At an 80% ratio, alerts drop to 274–619 alerts/hr at 30 FPS and 46–103 alerts/hr at 5 FPS. Hard-negative mining provides comparable suppression with one-third the dataset volume and one-sixth the negative frames: at 40%, hard-mined models produce 550–619 alerts/hr at 30 FPS, an 86% to 92% reduction from the 0% baseline depending on architecture.

These projections highlight an engineering distinction: *inference latency depends on architecture, whereas operational reliability depends on data curation*. While all four models execute within real-time limits (7.7–21.7 ms; Fig. 3), negative curation determines whether a detector avoids alert fatigue and thermal throttling.

V. DISCUSSION

A. Architectural Mechanisms and Practical Curation Guidelines

A detector’s feature-aggregation mechanism determines how it responds to negative-frame curation.

Mechanistic Divergence: Random negative sampling primarily selects low-entropy background patches that generate minimal training gradients. Purely convolutional networks such as YOLO26n rely on localized 3×3 re-parameterized convolutions (RepConv), which quickly learn repetitive cabin textures, causing YOLO26n to plateau at a 20% negative ratio and gain only modest benefit from hard-negative mining (–33.3%). Models using self-attention (YOLO11n), area attention (YOLO12n), or dual-label assignment (YOLOv10n) instead capture wider spatial context. Without negative frames, attention layers often mistake cabin lighting or seat contours for driver cues; mined hard negatives supply high-entropy distractors that force attention heads to distinguish cabin structures from driver actions. For YOLOv10n, mined negatives also penalize over-assignment in the auxiliary head, sharpening the primary branch’s decision boundary.

Per-Seed Stability: In our two-seed evaluations, attention-based and dual-assignment models showed consistent suppression across runs, whereas YOLO26n showed marked variability (0.0% under Seed 42 versus 47.0% under Seed 43). This suggests possible seed sensitivity in localized convolutional backbones, but a larger seed battery is needed before drawing definitive architectural conclusions.

Practical Curation Guidelines: These results establish two engineering guidelines for edge driver monitoring systems (DMS):

- 1) *Backbone-Targeted Curation:* Uniform random sampling suffices for convolutional backbones, but attention-centric and dual-assignment models need hard-negative mining, yielding an additional 52.6% to 67.8% false-alarm reduction at matched training compute.
- 2) *Prevalence Trade-Offs (40% vs. 80%):* An 80% negative ratio achieves the lowest false-alarm floor for most architectures (3.14–4.72 FP/1k), though YOLO12n reaches its minimum under 40% hard mining (6.29 FP/1k). Where storage or compute is constrained, a 40% hard-mined ratio approaches 80% random suppression (6.29–7.08 FP/1k) with only one-third the dataset volume and one-sixth the negative frames.

B. Limitations

This evaluation has six primary limitations:

Single-Annotator Labeling: A single annotator labeled the visual cues without independent review, which may introduce subjective bias in ambiguous edge cases.

Simplified Cue Annotation and Mutual Exclusion: Annotations enforced a single bounding box per positive frame and a strict mutual-exclusion rule between overlapping behaviors (e.g., a covered yawn is labeled exclusively as hand-over-mouth). In naturalistic driving, concurrent cues frequently co-occur; this protocol simplifies the detection objective and may understate real-world false negatives when simultaneous behaviors arise.

Single-Modality RGB Sensing: The evaluation relies entirely on visible-light RGB video without near-infrared (NIR) or active depth sensing commonly deployed for in-cabin night vision.

Isolated Frame-Level Inference: Detections were evaluated per frame without temporal persistence filtering (e.g., Kalman tracking or sliding time-windows), yielding an unmitigated upper-bound alert baseline.

Driver Partition Overlap and Single-Pass Mining: Splits were stratified across cue instances rather than subject-disjoint driver identities. Extracting 1 FPS frames from continuous sessions risks cross-split background leakage, framing reported suppression as an empirical upper bound until verified on subject-disjoint splits. Hard-negative mining was also restricted to a single closed-loop iteration rather than multi-round bootstrapping.

Desktop-GPU Deployment Proxy: Latency and memory profiling (Section III-F) was performed on a desktop-class discrete GPU rather than an automotive-grade embedded SoC, and no power draw was measured. Absolute latency, memory, and thermal margins on production in-cabin hardware may differ from the desktop-proxy figures reported here.

VI. CONCLUSION

This study evaluated negative-frame curation and hard-negative mining for suppressing false alarms in edge-deployed Driver Monitoring Systems (DMS). Across four You Only Look Once (YOLO) architectures under 3 million parameters, increasing negative training frames from 0% to 80%

reduced false positives per 1,000 frames (FP/1k) by $6.4\times$ to $25.7\times$. Core detection sensitivity (mAP_{50:95}) and recall stayed stable while precision increased. At matched dataset volume, replacing random negative frames with hard-mined false positives achieved a further 33.3% to 67.8% false-alarm reduction. Attention-based and dual-assignment models benefited most from hard negatives, whereas purely convolutional models achieved effective suppression from random background frames alone. In streaming projections, these curation protocols lowered nuisance alerts below 620 alerts/hr at 30 frames per second (FPS), mitigating driver alert fatigue and edge computational load.

These findings show that dataset curation is as critical as architectural design for edge deployment, and that curation strategy must match the detector's internal feature-aggregation mechanism: attention heads need high-entropy mined distractors to eliminate false correlations, whereas convolutional kernels regularize effectively on uniform background scenes.

This study has six primary shortcomings: visual cues were labeled by a single annotator; annotations enforced single-cue mutual exclusion rather than simultaneous multi-cue interactions; evaluation relied exclusively on single-frame visible-light RGB without temporal smoothing; mining was limited to a single pass; data splits were stratified across cue instances rather than disjoint drivers; and latency/memory profiling relied on a desktop-GPU proxy rather than validated automotive SoC hardware, with no power measurement. Future work should validate these protocols on multi-modal near-infrared (NIR) and depth sensors, incorporate multi-frame temporal consensus, evaluate iterative multi-round mining, test cross-subject generalization on disjoint driver splits, and repeat the latency and power-draw benchmarks on representative automotive SoC hardware.

REPRODUCIBILITY STATEMENT

To facilitate independent verification, training configuration scripts, hyperparameter files, and exact frame split indices will be made publicly available in an open-source repository upon paper acceptance.

REFERENCES

- [1] European Parliament and Council of the European Union, Regulation (EU) 2019/2144 on type-approval requirements for the general safety of motor vehicles, Art. 6 and Annex II (Advanced Driver Distraction Warning), 2019, as implemented by Commission Delegated Regulation (EU) 2023/2590 laying down detailed test procedures and technical requirements, 2023. Advanced Driver Distraction Warning (ADDW) systems are mandatory for newly type-approved vehicles from 7 July 2024 and for all newly registered vehicles from 7 July 2026.
- [2] J. D. Ortega, N. Kose, P. Cañas, M.-A. Chao, A. Unnervik, M. Nieto, O. Otaegui, and L. Salgado, "DMD: A large-scale multi-modal driver monitoring dataset for attention and alertness analysis," in *Proc. Eur. Conf. Comput. Vis. (ECCV) Workshops*, 2020, pp. 387–405.
- [3] R. Florez, F. Palomino-Quispe, A. B. Alvarez, R. J. Coaquira-Castillo, and J. C. Herrera-Levano, "A real-time embedded system for driver drowsiness detection based on visual analysis of the eyes and mouth using convolutional neural network and mouth aspect ratio," *Sensors*, vol. 24, no. 19, Art. no. 6261, Sep. 2024, doi: 10.3390/s24196261.
- [4] S. Liu, L. Liu, J. Tang, B. Yu, Y. Wang, and W. Shi, "Edge computing for autonomous driving: Opportunities and challenges," *Proc. IEEE*, vol. 107, no. 8, pp. 1697–1716, Aug. 2019, doi: 10.1109/JPROC.2019.2915983.

- [5] F. Lambert, "Tesla's self-driving safeguards fooled by \$30 doll heads," *Electrek*, Jun. 15, 2026. [Online]. Available: <https://electrek.co/2026/06/15/chinese-drivers-plastic-heads-fool-tesla-autopilot-camera/>.
- [6] T. Benoit-Cattin, D. Velasco-Montero, and J. Fernández-Berni, "Impact of thermal throttling on long-term visual inference in a CPU-based edge device," *Electronics*, vol. 9, no. 12, Art. no. 2106, 2020, doi: 10.3390/electronics9122106.
- [7] M. Martin, A. Roitberg, M. Haurilet, M. Horne, S. Reiß, M. Peng, and R. Stiefelhagen, "Drive&Act: A multi-modal dataset for fine-grained driver behavior recognition in autonomous vehicles," in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2019, pp. 2801–2810.
- [8] P. Cañas, J. D. Ortega, M. Nieto, and O. Otaegui, "Detection of distraction-related actions on DMD: An image and a video-based approach comparison," in *Proc. Int. Conf. Comput. Vis. Theory Appl. (VISAPP)*, 2021, pp. 458–465.
- [9] Z. Yu, J. Yu, S. Fong, and Y. Wu, "YOLO-Drive: Robust driver distraction recognition under fine-grained and overlapping behaviors," *Comput. Mater. Contin.*, vol. 87, no. 2, 2026, doi: 10.32604/cmc.2025.074899.
- [10] M. Chaman, A. El Maliki, H. El Yanboiy, H. Dahou, and A. Hadjoudja, "YOLOv12-based driver monitoring system: Real-time applications in ADAS," *Int. J. Saf. Secur. Eng.*, vol. 15, no. 9, pp. 1921–1932, 2025, doi: 10.18280/ijss.150915.
- [11] M. R. Elshamy, H. M. Emara, M. R. Shoaib, and A.-H. A. Badawy, "P-YOLOv8: Efficient and accurate real-time detection of distracted driving," *arXiv preprint arXiv:2410.15602*, 2024.
- [12] D. Herath, C. Abeyrathne, and P. Jayaweera, "Vision-based driver drowsiness monitoring: Comparative analysis of YOLOv5-v11 models," *arXiv preprint arXiv:2509.17498*, 2025.
- [13] R. Parasuraman and V. Riley, "Humans and automation: Use, misuse, disuse, abuse," *Hum. Factors*, vol. 39, no. 2, pp. 230–253, 1997.
- [14] A. Shrivastava, A. Gupta, and R. Girshick, "Training region-based object detectors with online hard example mining," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 761–769.
- [15] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, "Focal loss for dense object detection," in *Proc. IEEE Int. Conf. Comput. Vis. (ICCV)*, 2017, pp. 2980–2988.
- [16] G. Jocher, A. Chaurasia, and J. Qiu, "Ultralytics YOLO: Tips for best training results and background image curation," *Ultralytics Documentation*, 2023. [Online]. Available: <https://docs.ultralytics.com/yolov5/tutorials/tips-for-best-training-results/>
- [17] S. Jin, A. RoyChowdhury, H. Jiang, A. Singh, A. Prasad, D. Chakraborty, and E. Learned-Miller, "Unsupervised hard example mining from videos for improved object detection," in *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 2018, pp. 153–169.
- [18] J. S. Stutts, J. Feaganes, D. Reinfurt, E. Rodgman, C. Hamlett, K. Gish, and L. Staplin, "Driver's exposure to distractions in their natural driving environment," *Accid. Anal. Prev.*, vol. 37, no. 6, pp. 1093–1101, 2005, doi: 10.1016/j.aap.2005.06.007.
- [19] A. Wang, H. Chen, L. Liu, K. Chen, Z. Lin, J. Han, and G. Ding, "YOLOv10: Real-time end-to-end object detection," *arXiv preprint arXiv:2405.14458*, 2024.
- [20] R. Khanam and M. Hussain, "YOLOv11: An overview of the key architectural enhancements," *arXiv preprint arXiv:2410.17725*, 2024.
- [21] Y. Tian, Q. Ye, and D. Doermann, "YOLOv12: Attention-centric real-time object detectors," *arXiv preprint arXiv:2502.12524*, 2025.
- [22] G. Jocher, J. Qiu, M. Liu, S. Lyu, F. C. Akyon, and M. E. Kalfaoglu, "Ultralytics YOLO26: Unified real-time end-to-end vision models," *arXiv preprint arXiv:2606.03748*, 2026.